

SBA (Small Biz Admin.) Loan Approval Analysis & Prediction

by Jungseok Lee

```
In [1]: from IPython.display import display, HTML
display(HTML("<style>.container { width:80% !important; }</style>"))

import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import shap
import h2o
from h2o.estimators import H2OTargetEncoderEstimator
from h2o.estimators import H2OGradientBoostingEstimator
try:
    h2o.cluster().shutdown()
except:
    pass
```

```
In [2]: h2o.init(max_mem_size=8)
```

```
Checking whether there is an H2O instance running at http://localhost:54321..... not found.
Attempting to start a local H2O server...
; Java HotSpot(TM) 64-Bit Server VM (build 21.0.2+13-LTS-58, mixed mode, sharing)
Starting server from C:\Users\wizdo\ml-spring-2024\Lib\site-packages\h2o\backend\bin\h2o.jar
Ice root: C:\Users\wizdo\AppData\Local\Temp\tmpjnyajql
JVM stdout: C:\Users\wizdo\AppData\Local\Temp\tmpjnyajql\h2o_wizdo_started_from_python.out
JVM stderr: C:\Users\wizdo\AppData\Local\Temp\tmpjnyajql\h2o_wizdo_started_from_python.err
Server is running at http://127.0.0.1:54321
Connecting to H2O server at http://127.0.0.1:54321 ... successful.
Warning: Your H2O cluster version is (4 months and 8 days) old. There may be a newer version available.
Please download and install the latest version from: https://h2o-release.s3.amazonaws.com/h2o/latest_stable.html
```


	index	City	State	Zip	Bank	BankState	NAICS	NoEmp	
type	int	enum	enum	int	enum	enum	int	int	
mins	0.0			0.0			0.0	0.0	
mean	399677.5			53800.14742367606			398464.7250036278	11.394357457753493	1.28043
maxs	799355.0			99999.0			928120.0	9999.0	
sigma	230754.34522669343			31185.719098605336			263323.9798013513	73.98731938246111	0.45180
zeros	1			249			179717	5914	
missing	0	25	12	0	1402	1408	0	0	
0	0.0	FORT LEE	NJ	7024.0	BNB HANA BANK NATL ASSOC	NJ	425120.0	2.0	
1	1.0	WESTWEGO	LA	70094.0	JEDCO DEVELOPMENT CORPORATION	LA	812331.0	62.0	
2	2.0	DENVER	CO	80209.0	WELLS FARGO BANK NATL ASSOC	SD	541611.0	4.0	
3	3.0	WRANGELL	AK	99929.0	FIRST BANK	AK	446110.0	3.0	
4	4.0	MALVERN	AR	72104.0	CITICAPITAL SMALL BUS. FINANCE	TX	0.0	1.0	
5	5.0	HAMILTON SQUARE	NJ	8619.0	SUN NATIONAL BANK	NJ	445110.0	3.0	
6	6.0	PHOENIX	AZ	85040.0	MUTUAL OF OMAHA BANK	AZ	0.0	1.0	
7	7.0	DENVER	CO	80207.0	JPMORGAN CHASE BANK NATL ASSOC	IL	722211.0	11.0	
8	8.0	CLEVELAND	OH	44109.0	U.S. BANK NATIONAL ASSOCIATION	OH	445299.0	4.0	
9	9.0	ESCONDIDO	CA	92025.0	COMERICA BANK	TX	621210.0	4.0	

Clean up (Encode replace missing values)

```
In [5]: columns = ["City", "State", "Bank", "BankState",
                  "NewExist", "RevLineCr", "LowDoc"]
num_col = data.columns_by_type(coltype="numeric")
enum_col = data.columns_by_type(coltype="categorical")
all_columns = data.columns

for col in columns:
    if all_columns.index(col)*1.0 in num_col:
        print("Fillna for numerical column:...", col)
        '''Alternative way '''
        data[data[col].isna(), col] = 0
    elif all_columns.index(col)*1.0 in enum_col:
        print("Fillna for categorical column:...", col)
        data[col] = data[col].ascharacter()
        data[data[col].isna(), col] = "Missing"
        data[col] = data[col].asfactor()
```

```
Fillna for categorical column:... City
Fillna for categorical column:... State
Fillna for categorical column:... Bank
Fillna for categorical column:... BankState
Fillna for numerical column:... NewExist
Fillna for categorical column:... RevLineCr
Fillna for categorical column:... LowDoc
```

```
In [ ]: # missing values were filled with 0 or 'Missing'
```

Split dataset to Train/Valid/Test

```
In [6]: train, valid, test = data.split_frame(ratios=[.6, .2], seed=1234)
```

Change to categorical variables for feature engineering

```
In [7]: cat_columns = ["City", "State", "Bank", "BankState", "UrbanRural", "FranchiseCode",
                      "NewExist", "RevLineCr", "LowDoc", "Zip", "NAICS"]
```

```

encoded_columns = cat_columns
response = "MIS_Status"

train[encoded_columns+[response]] = train[encoded_columns+[response]].asfactor()
valid[encoded_columns+[response]] = valid[encoded_columns+[response]].asfactor()
test[encoded_columns+[response]] = test[encoded_columns+[response]].asfactor()

```

Add engineered features

(Excluded some features that didn't improve the model performance)

```
In [8]: train_0 = train
```

```
In [9]: import numpy as np
```

```

column_to_bin = col = "BalanceGross"
col_np_array = np.array(train_0[column_to_bin].as_data_frame(use_pandas=True)[column_to_bin].values, dtype=np.int64) # Convert sin
counts, breaks = np.histogram(col_np_array, bins=5)

```

converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).

```

In [10]: min_val = min(col_np_array)-1                                # Establish min and max values
          max_val = max(col_np_array)+1

          new_b = [int(min_val)]                                       # Redefine breaks such that each bucket has enough support
          for i in range(4):
              if counts[i] > 1000 and counts[i+1] > 1000:
                  new_b.append(int(breaks[i+1]))
          new_b.append(int(max_val))

          names = [col + '_' + str(x) for x in range(len(new_b)-1)] # Generate names for buckets, these will be categorical names

```

```
In [11]: train[col+"_cut"] = train[col].cut(breaks=new_b, labels=names)
```

```

In [12]: def cut_column(col_np_array, train, valid, test, col):

          counts, breaks = np.histogram(col_np_array, bins=100)      # Generate counts and breaks for our histogram
          min_val = min(col_np_array)-1                               # Establish min and max values
          max_val = max(col_np_array)+1

          new_b = [int(min_val)]                                       # Redefine breaks such that each bucket has enough support
          for i in range(99):

```

```

        if counts[i] > 1000 and counts[i+1] > 1000:
            new_b.append(int(breaks[i+1]))
    new_b.append(int(max_val))

    names = [col + '_' + str(x) for x in range(len(new_b)-1)] # Generate names for buckets, these will be categorical names

    train[col+"_cut"] = train[col].cut(breaks=new_b, labels=names)
    valid[col+"_cut"] = valid[col].cut(breaks=new_b, labels=names)
    test[col+"_cut"] = test[col].cut(breaks=new_b, labels=names)

```

```

In [13]: def add_features(train, valid, test):

    """
    Helper function to add binning and interaction features to the dataset
    """

    # Transform numerical columns to categorical via binning
    for column_to_bin in ["NoEmp",
                          "CreateJob",
                          "RetainedJob",
                          "DisbursementGross",
                          "BalanceGross",
                          "GrAppv",
                          "SBA_Appv"]:
        print("Binning Column: {}".format(column_to_bin))
        col_np_array = np.array(train[column_to_bin].as_data_frame(use_pandas=True)[column_to_bin].values, dtype=np.int64) # Convert to numpy array
        cut_column(cut_np_array, train, valid, test, column_to_bin)

    #Add interaction columns for a subset of columns
    interaction_cols1 = ["NoEmp_cut",
                        "City",
                        "State",
                        "Bank",
                        "BankState",
                        "NAICS",
                        "NewExist",
                        "CreateJob_cut",
                        "RetainedJob_cut",
                        "UrbanRural",
                        "RevLineCr",
                        "LowDoc",
                        "FranchiseCode",
                        "DisbursementGross_cut",
                        "BalanceGross_cut",
                        "GrAppv_cut",
                        "SBA_Appv_cut"]

```

```

train_cols = train.interaction(factors=interaction_cols1,      #Generate pairwise columns
                               pairwise=True,
                               max_factors=1000,
                               min_occurrence=100,
                               destination_frame="itrain")
valid_cols = valid.interaction(factors=interaction_cols1,
                               pairwise=True,
                               max_factors=1000,
                               min_occurrence=100,
                               destination_frame="invalid")
test_cols = test.interaction(factors=interaction_cols1,
                             pairwise=True,
                             max_factors=1000,
                             min_occurrence=100,
                             destination_frame="itest")

train = train.cbind(train_cols)      #Append pairwise columns to H2OFrames
valid = valid.cbind(valid_cols)
test = test.cbind(test_cols)

print("All columns added via interaction",train_cols.columns)

return train, valid, test

```

In [14]: train_f, valid_f, test_f = add_features(train, valid, test)

Binning Column: RetainedJob

converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).

Binning Column: BalanceGross

converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).

Binning Column: SBA_Appv

converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).


```
Out[16]: Model Details
=====
H2OTargetEncoderEstimator : TargetEncoder
Model Key: TargetEncoder_model_python_1714353241713_1
```

Target Encoder model summary: Summary for target encoder model

original_names	encoded_column_names
City	City_te
State	State_te
Bank	Bank_te
BankState	BankState_te
FranchiseCode	FranchiseCode_te
RevLineCr	RevLineCr_te
LowDoc	LowDoc_te

```
[tips]
Use `model.explain()` to inspect the model.
--
Use `h2o.display.toggle_user_tips()` to switch on/off this section.
```

Save target encoder path

```
In [17]: encoder_path = h2o.save_model(model=encoder_te, path="./artifacts", force=True)
```

```
In [18]: encoder_path
```

```
Out[18]: 'C:\\Users\\wizdo\\Downloads\\artifacts\\TargetEncoder_model_python_1714353241713_1'
```

Add the features created through target encoding.

```
In [19]: train_te = encoder_te.transform(frame=train_f, as_training=True)
valid_te = encoder_te.transform(frame=valid_f, as_training=False)
```

```
test_te = encoder_te.transform(frame=test_f, as_training=False)
```

Prepare for model training and tuning.

```
In [20]: predictors = train_te.columns
```

```
In [21]: for col in cat_columns+[response]:  
         predictors.remove(col)
```

Remove "index" column for training

```
In [22]: predictors.remove("index")
```

Model training and tuning (GBM(H2O))

```
In [23]: best_auc = 0  
best_nfolds = 0  
best_ntrees = 0  
best_max_depth = 0  
best_stopping_rounds = 0  
best_stopping_metric = []  
best_model = None  
  
for nfolds in [3, 4, 5]:  
    for ntrees in [30, 50, 80, 100, 130]:  
        for stopping_rounds in [3, 5]:  
            for max_depth in [3, 5]:  
                for stopping_metric in ["MAE"]:  
                    model = H2OGradientBoostingEstimator(nfolds = nfolds,  
                                                         ntrees = ntrees,  
                                                         max_depth = max_depth,  
                                                         stopping_rounds = stopping_rounds,  
                                                         stopping_metric = stopping_metric,  
                                                         seed=1234,  
                                                         keep_cross_validation_predictions = False)  
                    model.train(y=response, x=predictors, training_frame=train_te, validation_frame=valid_te)  
  
                    auc = model.model_performance(valid=True).auc()  
                    if auc > best_auc:  
                        best_auc = auc
```

```
best_nfolds = nfolds
best_ntrees = ntrees
best_max_depth = max_depth
best_stopping_rounds = stopping_rounds
best_stopping_metric = stopping_metric
best_model = model
print("Better Model found. Trained H2OGBM with (auc, nfolds, ntrees, stopping_rounds, stopping_metric):(",
      auc, nfolds, ntrees, max_depth, stopping_rounds, stopping_metric,")")

print("Best model AUC(on valid dataset):", best_auc)
print("Best model AUC(on test dataset):", best_model.model_performance(test_data=test_te).auc())
```

[illegible]

```
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
gbm Model Build progress: ██████████ (done) 100%
Best model AUC(on valid dataset): 0.8260729441934694
Best model AUC(on test dataset): 0.8231652178480686
```

Best model AUC(on test dataset): 0.8231652178480686

Final metrics using Test dataset

```
In [24]: best_model.model_performance(test_data=test_te)
```

Out[24]: ModelMetricsBinomial: gbm
** Reported on test data. **

MSE: 0.11013973593898531
RMSE: 0.33187307203053595
LogLoss: 0.355642847141753
Mean Per-Class Error: 0.27068745817623496
AUC: 0.8231652178480686
AUCPR: 0.5339127974547134
Gini: 0.6463304356961372

Confusion Matrix (Act/Pred) for max f1 @ threshold = 0.23895010495437083

	0	1	Error	Rate
0	111374.0	20312.0	0.1542	(20312.0/131686.0)
1	10780.0	17066.0	0.3871	(10780.0/27846.0)
Total	122154.0	37378.0	0.1949	(31092.0/159532.0)

Maximum Metrics: Maximum metrics at their respective thresholds

	metric	threshold	value	idx
	max f1	0.2389501	0.5233043	231.0
	max f2	0.1091954	0.6340495	307.0
	max f0point5	0.4486483	0.5410389	144.0
	max accuracy	0.4890380	0.8502683	130.0
	max precision	0.9741083	1.0	0.0
	max recall	0.0037266	1.0	399.0
	max specificity	0.9741083	1.0	0.0
	max absolute_mcc	0.3038338	0.4145142	202.0
	max min_per_class_accuracy	0.1543333	0.7422277	278.0
	max mean_per_class_accuracy	0.1453233	0.7436559	283.0

	metric	threshold	value	idx
	max tns	0.9741083	131686.0	0.0
	max fns	0.9741083	27829.0	0.0
	max fps	0.0037266	131686.0	399.0
	max tps	0.0037266	27846.0	399.0
	max tnr	0.9741083	1.0	0.0
	max fnr	0.9741083	0.9993895	0.0
	max fpr	0.0037266	1.0	399.0
	max tpr	0.0037266	1.0	399.0

Gains/Lift Table: Avg response rate: 17.45 %, avg score: 16.96 %

group	cumulative_data_fraction	lower_threshold	lift	cumulative_lift	response_rate	score	cumulative_response_rate	cumulative
1	0.0100043	0.8394239	4.8496171	4.8496171	0.8464912	0.8823035	0.8464912	0.88
2	0.0200023	0.7845444	4.3497916	4.5997827	0.7592476	0.8117718	0.8028831	0.84
3	0.0300003	0.7340418	4.0193367	4.4063411	0.7015674	0.7596531	0.7691183	0.81
4	0.0400045	0.6849076	3.7727221	4.2478867	0.6585213	0.7097847	0.7414604	0.79
5	0.0500025	0.6377627	3.6314115	4.1246226	0.6338558	0.6604728	0.7199448	0.76
6	0.1000050	0.4650787	2.9855574	3.5550900	0.5211232	0.5435609	0.6205340	0.65
7	0.1500013	0.3550272	2.2669234	3.1257370	0.3956871	0.4061945	0.5455913	0.57
8	0.2000038	0.2776231	1.8069912	2.7960402	0.3154068	0.3135642	0.4880434	0.50
9	0.3000025	0.1815220	1.3833399	2.3251500	0.2414593	0.2249186	0.4058504	0.41
10	0.4000013	0.1240623	0.9829183	1.9895973	0.1715665	0.1505517	0.3472803	0.34
11	0.5	0.0866626	0.7096261	1.7336063	0.1238639	0.1040273	0.3025976	0.29
12	0.5999987	0.0621321	0.5056445	1.5289481	0.0882593	0.0734980	0.2668749	0.26
13	0.6999975	0.0451910	0.3892888	1.3661411	0.0679496	0.0531865	0.2384573	0.23
14	0.7999962	0.0322940	0.2503084	1.2266631	0.0436908	0.0385311	0.2141117	0.20

group	cumulative_data_fraction	lower_threshold	lift	cumulative_lift	response_rate	score	cumulative_response_rate	cumulative
15	0.8999950	0.0203699	0.1436490	1.1063290	0.0250737	0.0264381	0.1931076	0.18
16	1.0	0.0020049	0.0430920	1.0	0.0075216	0.0111794	0.1745481	0.16

```
In [65]: best_model.model_performance(test_data=test_te).auc()
```

```
Out[65]: 0.8231652178480686
```

```
In [25]: best_model.model_performance(test_data=test_te).find_threshold_by_max_metric("f1")
```

```
Out[25]: 0.23895010495437083
```

```
In [26]: best_model.model_performance(test_data=test_te).confusion_matrix()
```

```
Out[26]: Confusion Matrix (Act/Pred) for max f1 @ threshold = 0.23895010495437083
```

	0	1	Error	Rate
0	111374.0	20312.0	0.1542	(20312.0/131686.0)
1	10780.0	17066.0	0.3871	(10780.0/27846.0)
Total	122154.0	37378.0	0.1949	(31092.0/159532.0)


The threshold 0.23895 is the boundary for predicting that the predicted value of 'MIS_status' is 0 or 1 when test_te data is entered into best_model. This threshold of 0.23895 is the value that maximizes the f1 value.


Shapley values

```
In [28]: def examine_all_reason_codes(data, model):

    shap_contribs = model.predict_contributions(data)

    col_mapping = {}
    for i in data.col_names:
        related_cols = [x for x in shap_contribs.col_names if "{}.".format(i) in x]
        if len(related_cols) > 0:
            col_mapping[i] = related_cols
```




Summary plot with Shapley values

```
In [30]: best_model.explain(test_te, include_explanations=['shap_summary']);
```

SHAP Summary

SHAP summary plot shows the contribution of the features for each instance (row of data). The sum of the feature contributions and the bias term is equal to the raw prediction of the model, i.e., prediction before applying inverse link function.


SHAP Summary plot for "GBM_model_python_1714353241713_2823"



Permutation feature importance

```
In [31]: best_model.permutation_importance_plot(test_te)
```

Permutation Variable Importance: gbm



Out[31]: Variable Importances

Variable	Relative Importance	Scaled Importance	Percentage
Bank_te	0.0284783	1.0	0.1864774
Bank_FranchiseCode	0.0192211	0.6749392	0.1258609
City_te	0.0169819	0.5963100	0.1111984
NAICS_FranchiseCode	0.0166022	0.5829771	0.1087121
Bank_RevLineCr	0.0122859	0.4314117	0.0804486
Bank_BankState	0.0121552	0.4268239	0.0795930
State_Bank	0.0047715	0.1675488	0.0312441
FranchiseCode_te	0.0045317	0.1591295	0.0296741
DisbursementGross	0.0039888	0.1400627	0.0261185
UrbanRural_FranchiseCode	0.0039447	0.1385150	0.0258299
---	---	---	---
BankState_BalanceGross_cut	0.0	0.0	0.0
BalanceGross_cut_SBA_Appv_cut	0.0	0.0	0.0
NewExist_BalanceGross_cut	0.0	0.0	0.0
BalanceGross_cut	0.0	0.0	0.0
UrbanRural_BalanceGross_cut	0.0	0.0	0.0
RevLineCr_LowDoc	0.0	0.0	0.0
BalanceGross	0.0	0.0	0.0
RetainedJob_cut_SBA_Appv_cut	0.0	0.0	0.0
RetainedJob_cut_LowDoc	0.0	0.0	0.0
NewExist_LowDoc	0.0	0.0	0.0

[87 rows x 4 columns]

What are the most important features, how they impact model predictions.

According to the Summary plot with Shapley values(or Shap summary), City_te, Bank_FranchiseCode, Bank_te are the most important features in that order. According to the Permutation feature importance graph, Bank_te, Bank_FranchiseCode, City_te are the most important features in that order.

Features affect the model prediction of 'MIS_status' by their shap values. For example, According to the Shap summary, the higher the 'City_te', the higher the shap value. In other words, Higher 'City_te' value has an influence in the direction of predicting that 'MIS_status' is 1.

Individual observations analysis using Shapley values.

```
In [39]: test_te["MIS_Status"].head(20)
```

Out[39]:

MIS_Status

0

1

0

0

0

1

0

0

0

0

0

0

0

0

1

0

0

0

0

1

[20 rows x 1 column]

In [40]:

```
best_model.predict(test_te).head(20)
```

gbm prediction progress:

(done) 100%

Out[40]:

	predict	p0	p1
	0	0.979175	0.0208253
	0	0.952072	0.0479278
	1	0.581554	0.418446
	0	0.896618	0.103382
	0	0.912671	0.0873286
	0	0.906274	0.0937263
	0	0.959727	0.0402734
	1	0.66142	0.33858
	0	0.96091	0.0390896
	0	0.897612	0.102388
	0	0.993964	0.00603563
	1	0.573491	0.426509
	0	0.986225	0.0137746
	1	0.527061	0.472939
	1	0.234944	0.765056
	0	0.971343	0.0286574
	0	0.973032	0.0269675
	0	0.981235	0.018765
	0	0.84971	0.15029
	1	0.595205	0.404795

[20 rows x 3 columns]

Compare the 'MIS_Status' value of the test data and the predicted value obtained by inserting the test data into best_model, and check the correctly predicted observations(row) and incorrectly predicted observations(row) for Individual observations analysis using Shapley values.

```
In [41]: # The function is to calculate Shapley values (contributions) and plot them for single record
def examine_indiv_reason_codes(record, model, use_matplotlib=True):
```

```
shap_contribs = model.predict_contributions(record)
```

```
col_mapping = {}
for i in record.col_names:
    related_cols = [x for x in shap_contribs.col_names if "{}.".format(i) in x]
    if len(related_cols) > 0:
        col_mapping[i] = related_cols

for k, v in col_mapping.items():
    if len(v) > 1:
        shap_contribs[k] = shap_contribs[v].sum(axis=1, return_frame=True)
        shap_contribs = shap_contribs.drop(v)

shap_cols = [i for i in shap_contribs.col_names if i != "BiasTerm"]
bias_term = shap_contribs.as_data_frame()["BiasTerm"].values
X = record.as_data_frame(use_pandas=True)
shap_contribs = shap_contribs.as_data_frame(use_pandas=True)

return shap.force_plot(bias_term,
                        shap_contribs[shap_cols].values,
                        X[shap_cols].values,
                        shap_cols,
                        link="logit",
                        matplotlib=use_matplotlib
                    )
```

```
In [42]: shap.initjs()
```



Label 0 is correctly identified


```
In [43]: examine indiv reason codes(test te[0, :], best model, use matplotlib=False)
```

[illegible]

```

converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).

```




The p1 value of row0 observation is 0.02. 'State_Bank' has the most impact in the direction of predicting that 'MIS_status' is 1. 'City_te' has the most impact in the direction of predicting that 'MIS_status' is 0.

```
In [44]: examine_indiv_reason_codes(test_te[3, :], best_model, use_matplotlib=False)
```

```
contributions progress: ██████████ | (done) 100%
```

```
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
```

Out[44]:



The p1 value of row3 observation is 0.10. 'City_te' has the most impact in the direction of predicting that 'MIS_status' is 1. 'NAICS_FranchiseCode' has the most impact in the direction of predicting that 'MIS_status' is 0.

Label 0 is identified as 1

```
In [45]: examine_indiv_reason_codes(test_te[2, :], best_model, use_matplotlib=False)
```

```
contributions progress: |██████████████████████████████████████████| (done) 100%
```

```
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
```

Out[45]:



The p1 value of row2 observation is 0.42. 'Bank_FranchiseCode' has the most impact in the direction of predicting that 'MIS_status' is 1. 'NAICS_FranchiseCode' has the most impact in the direction of predicting that 'MIS_status' is 0.

```
In [46]: examine_indiv_reason_codes(test_te[7, :], best_model, use_matplotlib=False)
```


[illegible]

```

converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).

```

Out[46]:



The p1 value of row7 observation is 0.34. 'Bank_RevLineCr' has the most impact in the direction of predicting that 'MIS_status' is 1. 'Bank_FranchiseCode' has the most impact in the direction of predicting that 'MIS_status' is 0.

Label 1 is correctly identified

```
In [48]: examine_indiv_reason_codes(test_te[14, :], best_model, use_matplotlib=False)
```

```
contributions progress: ████████████████████████████████████████| (done) 100%
```

```
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
```

base value

0.0007028 0.001908 0.00517 0.01393 0.03698 0.09451 0.221

BankState_RevLineCr = NC_N Bank_te = 0.2814 Bank_RevLineCr = BANK OF AMERICA NATL ASSOC_N Bank_FranchiseCode =

```
examine_indiv_reason_codes(test_te[19, :], best_model, use_matplotlib=False)
```

Segment	Value
ZENS BANK NATL ASSOC_N	0.00517
NAICS_NewExist = 238990_1	0.01393
NAICS_SBA_Appv_cut = 238990_SBA_Appv_0	0.03698
Bank_te = 0.2236	0.09451
Bank_FranchiseCode = CITIZENS BA	0.221

```
In [50]: examine_indiv_reason_codes(test_te[1, :], best_model, use_matplotlib=False)
```

```
contributions progress: |██████████████████████████████████████████████████| (done) 100%  
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).  
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).  
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
```

Diagram illustrating the concept of a "base value" and its relationship to a function $f(x)$.

The horizontal axis represents the function $f(x)$, with values ranging from 0.008495 to 0.1468. The "base value" is marked at 0.05.

The diagram shows two regions relative to the base value:

- Higher Region (Red):** Values greater than the base value (0.05). The region is labeled "higher" and "f(x)".
- Lower Region (Blue):** Values less than the base value (0.05). The region is labeled "lower".

Key values on the axis include 0.008495, 0.01393, 0.02276, 0.03698, 0.05, 0.05954, 0.09451, and 0.1468.

Below the axis, a pink bar represents the "Rural = COMERICA BANK_1" region. A red box highlights the region between 0.01393 and 0.03698, labeled "Bank_BankState = COMERICA BANK_TX". A blue box highlights the region between 0.05954 and 0.09451, labeled "NAICS_RevLineCr = 621210_0". A blue box highlights the region between 0.09451 and 0.1468, labeled "NAICS_FranchiseCode = 621210_1". A blue box highlights the region between 0.1468 and 0.1468, labeled "NAICS".

```
In [51]: examine_indiv_reason_codes(test_te[5, :], best_model, use_matplotlib=False)
```

```
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
converting H2O frame to pandas dataframe using single-thread. For faster conversion using multi-thread, install datatable (for Python 3.9 or lower), or polars and pyarrow (for Python 3.10 or above).
```

higher ⇌ lower
base value

0.02276 0.03698 0.05954 0.0951 0.1468

kState = BANK OF AMERICA NATL ASSOC_NC Bank_UrbanRural = BANK OF AMERICA NATL ASSOC_1 Bank_te = 0.2716 City_te = 0.1352 Bank_RevLineC

Residuals analysis identify and report common patterns in the errors made by the model

```
In [52]: predict_MIS = best_model.predict(test_te)["predict"]
predict_MIS.columns = ["predict_MIS"]
test_te_p = test_te.cbind(predict_MIS)
test_te_p[["MIS Status", "predict MIS"]] = test_te_p[["MIS Status", "predict MIS"]].asnumeric()
```



```
Out[56]:
Model Details
=====
H2OGradientBoostingEstimator : Gradient Boosting Machine
Model Key: GBM_model_python_1714353241713_12885
```

Model Summary:

number_of_trees	number_of_internal_trees	model_size_in_bytes	min_depth	max_depth	mean_depth	min_leaves	max_leaves	mean_leave
130.0	130.0	175900.0	5.0	5.0	5.0	22.0	32.0	30.72307



ModelMetricsRegression: gbm
** Reported on train data. **

MSE: 0.08730088487242944
RMSE: 0.2954672314698018
MAE: 0.2030319585492246
RMSLE: 0.20606227138792627
Mean Residual Deviance: 0.08730088487242944

ModelMetricsRegression: gbm
** Reported on cross-validation data. **

MSE: 0.1223206358500482
RMSE: 0.3497436716368835
MAE: 0.24192716919818685
RMSLE: 0.24607900200183777
Mean Residual Deviance: 0.1223206358500482

Cross-Validation Metrics Summary:

	mean	sd	cv_1_valid	cv_2_valid	cv_3_valid
mae	0.2419228	0.0007219	0.2425398	0.2411289	0.2420997
mean_residual_deviance	0.1223159	0.0007594	0.1230415	0.1215267	0.1223795
mse	0.1223159	0.0007594	0.1230415	0.1215267	0.1223795

	mean	sd	cv_1_valid	cv_2_valid	cv_3_valid
r2	0.1751850	0.0021526	0.1730136	0.1773184	0.1752231
residual_deviance	0.1223159	0.0007594	0.1230415	0.1215267	0.1223795
rmse	0.3497358	0.0010859	0.3507727	0.3486068	0.3498279
rmsle	0.2460746	0.0006137	0.2467145	0.2454908	0.2460186

Scoring History:


timestamp	duration	number_of_trees	training_rmse	training_mae	training_deviance
2024-04-28 22:19:37	27.388 sec	0.0	0.3850945	0.2965956	0.1482978
2024-04-28 22:19:37	27.474 sec	1.0	0.3791827	0.2918395	0.1437795
2024-04-28 22:19:37	27.536 sec	2.0	0.3739977	0.2873904	0.1398743
2024-04-28 22:19:37	27.594 sec	3.0	0.3695775	0.2833019	0.1365876
2024-04-28 22:19:37	27.652 sec	4.0	0.3657425	0.2795468	0.1337676
2024-04-28 22:19:37	27.709 sec	5.0	0.3624543	0.2761083	0.1313731
2024-04-28 22:19:37	27.767 sec	6.0	0.3595890	0.2729529	0.1293043
2024-04-28 22:19:37	27.824 sec	7.0	0.3566522	0.2697047	0.1272008
2024-04-28 22:19:37	27.895 sec	8.0	0.3542547	0.2668747	0.1254964
2024-04-28 22:19:37	27.958 sec	9.0	0.3522416	0.2643809	0.1240741
---	---	---	---	---	---
2024-04-28 22:19:40	30.938 sec	58.0	0.3173201	0.2213509	0.1006920
2024-04-28 22:19:40	30.995 sec	59.0	0.3168681	0.2209577	0.1004054
2024-04-28 22:19:40	31.051 sec	60.0	0.3163780	0.2205496	0.1000950
2024-04-28 22:19:40	31.106 sec	61.0	0.3161196	0.2203247	0.0999316
2024-04-28 22:19:40	31.159 sec	62.0	0.3157734	0.2200161	0.0997128
2024-04-28 22:19:40	31.222 sec	63.0	0.3152301	0.2195304	0.0993700
2024-04-28 22:19:40	31.276 sec	64.0	0.3149679	0.2192934	0.0992048

timestamp	duration	number_of_trees	training_rmse	training_mae	training_deviance
2024-04-28 22:19:40	31.332 sec	65.0	0.3145582	0.2189551	0.0989469
2024-04-28 22:19:41	31.386 sec	66.0	0.3143027	0.2187589	0.0987862
2024-04-28 22:19:44	34.493 sec	130.0	0.2954672	0.2030320	0.0873009


[68 rows x 7 columns]

Variable Importances:

variable	relative_importance	scaled_importance	percentage
Bank_FranchiseCode	6952.7011719	1.0	0.1357536
NAICS_FranchiseCode	3892.1289062	0.5598010	0.0759950
State_Bank	3685.3530273	0.5300606	0.0719577
State_SBA_Appv_cut	3656.4809570	0.5259080	0.0713939
NAICS_UrbanRural	3463.7934570	0.4981939	0.0676316
NAICS_NewExist	2738.8505859	0.3939261	0.0534769
NAICS_RevLineCr	2468.5207520	0.3550449	0.0481986
State_BankState	2378.6149902	0.3421138	0.0464432
Bank_RevLineCr	2177.3486328	0.3131659	0.0425134
Bank_te	2018.9973145	0.2903903	0.0394215
---	---	---	---
UrbanRural_RevLineCr	0.0	0.0	0.0
UrbanRural_LowDoc	0.0	0.0	0.0
UrbanRural_BalanceGross_cut	0.0	0.0	0.0
RevLineCr_LowDoc	0.0	0.0	0.0
RevLineCr_BalanceGross_cut	0.0	0.0	0.0
LowDoc_FranchiseCode	0.0	0.0	0.0
LowDoc_BalanceGross_cut	0.0	0.0	0.0
LowDoc_SBA_Appv_cut	0.0	0.0	0.0




Bank_UrbanRural



```
In [58]: res_model.explain(test_te_p_abs_sort.tail(10000), include_explanations = ['shap_summary']);
```

SHAP Summary

SHAP summary plot shows the contribution of the features for each instance (row of data). The sum of the feature contributions and the bias term is equal to the raw prediction of the model, i.e., prediction before applying inverse link function.



According to the residuals analysis, `abs_error`(absolute value of the difference between 'MIS_Status' value and predicted value) made by the model shows patterns like above Shap summary. In other words, Bank_FranchiseCode, City_te, Bank_te have the greatest influence on the occurrence of `abs_error` in that order.

Save GBM(H2O) model path

```
In [59]: model_path = h2o.save_model(model=best_model, path="./artifacts", force=True)
```

```
In [60]: model_path
```

```
Out[60]: 'C:\\Users\\wizdo\\Downloads\\artifacts\\GBM_model_python_1714353241713_2823'
```

Summary and Conclusion

In this project, I created and verified a model that predicts MIS_Status values using SBA_loan data. Let me explain in order. First, I loaded the data and checked the characteristics of the data, such as the type of features and missing values. And the missing values were filled with 0 or 'Missing'. Next, the dataset was divided into train, valid, and test in a 6:2:2 ratio. And as a preliminary work of feature engineering, some variables were changed to categorical variables. After that I added engineered features. Additionally, 7 variables were converted through target encoding and this target encoder was saved for the scoring funct value.

Next, I removed the index and did model training and tuning. In this process, the most appropriate parameters were found and the best model was determined. As a result of 60 combinations, the auc value of the valid data was the highest(0.82607) under the condition (n folds = 3, ntrees = 130, max_depth = 5, stopping_rounds = 3, stopping_metric = MAE), so this model was set as the best model.

By inserting the test dataset into this best model, I obtained the auc value(auc: 0.82316), the threshold value that maximize f1(threshold = 0.23895), and the confusion matrix.

Likewise, I drew a Shap plot, Shap summary, and permutation feature importance graph using this best model and test dataset. Through this, the most important features(City_te, Bank_FranchiseCode, Bank_te) and the direction of their influence were identified. In addition, eight sample graphs were drawn to illustrate how much each feature affected the prediction when the the predicted value of each row matched MIS_Status value or not.

Next, I analyzed the residuals (the absolute value of the difference between MIS_Status and the predicted value) resulting from this model. As a result, I drew a Shap plot and Shap Summary graph and identified the features(Bank_FranchiseCode, City_te, Bank_te) that most affect these errors.

Finally, the best_model was saved for the scoring function and the storage location was checked.

To further develop the results of this project, I would like to suggest two more things. First, the auc value in the train dataset was around 0.92, but in the valid and test datasets, both came out to only around 0.82. If this result is not a characteristic of the dataset but a result of the model overfitting the train dataset, additional work is required. Second, some numerical features of this dataset were positively skewed. But I didn't utilize minmax scaler or log scaler. If this process had been added, model performance would have been better.

```
In [ ]:
```